

Dhruv Sharma

408-784-8712 | dvsharma@ucdavis.edu | in/dhruvvsharma1/ | github.com/dsharma136

EDUCATION

Georgia Institute of Technology

M.S. Computer Science, Part-Time (Online)

January 2027 - December 2028

University of California Davis

Bachelor of Science in Computer Science

December 2026

EXPERIENCE

AI Associate Consultant Intern | IBM Expert Labs

May 2026 - August 2026

- Built a knowledge agent for an enterprise government client in watsonx Orchestrate, defining a 150-question benchmark that quantified a 25% token expenditure reduction from instruction-set optimization.
- Drove the roadmap for an internal agent evaluation platform used by 30+ consultants at IBM, running stakeholder demos and converting feedback into a prioritized backlog shipping 5 features.
- Built the platform's evaluation and observability pipeline: LLM-as-judge scoring, human annotation with judge-alignment feedback, and per-question tool-call tracing and latency monitoring; parallelized execution cut runtime 4.4x.

Business AI Adoption Intern | SAP

September 2025 - March 2026

- Analyzed 10K+ rows of SAP BTP adoption data across Business AI products (including Joule) using Python and Excel to identify drop-off patterns and cross-market-unit performance gaps.
- Contributed analysis and slides to bi-weekly executive presentations covering 6 Americas market units, translating adoption metrics into insights that informed leadership discussions on customer engagement.
- Supported customer discovery calls with 5+ enterprise accounts, capturing AI adoption friction points and synthesizing findings into product feedback for the BTP team.

Product Management Intern | Nutanix

May 2025 - August 2025

- Conducted competitive research on Nutanix Unified Storage, Kubernetes Platform, and Enterprise AI, benchmarking AI workloads to identify technical and GTM gaps.
- Collaborated with PM, engineering, marketing, sales, and stakeholders to define semantic-search and metadata-indexing requirements and architecture, improving data discoverability by 60%.
- Developed a 6-month execution plan and multi-year roadmap for intelligent storage with semantic search, and built a PoC prototype projecting a 40% reduction in query latency.

Data Science Intern | Stanford Gevaert Lab

May 2023 - December 2024

- Optimized a deep learning model for microvascular segmentation on 1,622 diverse Kaggle medical images using Stanford's HPC cluster, achieving a dice score of 0.77.
- Trained a skin segmentation model for 150 epochs on 45,000+ images, leveraging CPU/GPU resources to achieve a 0.987 dice score and a 10% accuracy boost through refined masks.

Lead Machine Learning Intern | Idaho National Laboratory

July 2024 - September 2024

- Developed a LLaMA-based model trained on approximately 88 equipment qualification procedures corresponding to 933 drawings, streamlining documentation processes.
- Implemented OCR models to extract text from CAD drawings, achieving approximately 85% accuracy in providing structured data for machine learning tasks.

PROJECTS

TechBytes (Micro-Learning Application)

- Built a card-based web app delivering approximately 50 daily AI-summarized tech news blurbs across 6 categories, architected with a LangGraph orchestration pipeline scraping HackerNews, ArXiv, and Brave Search APIs and using Groq's Llama 3.3 70B API for inference, served via a React frontend.
- Researched competitors including TechCrunch and Deepstash to find whitespace for a tech-exclusive micro-learning product, positioning TechBytes as the first card-based daily learning experience for tech news.

TECHNICAL SKILLS

Languages: Python, JavaScript, Java, C/C++, R, Bash/Shell Scripting, SQL

AI/ML: LangGraph, LangChain, RAG, LLMs, Vector Databases (Milvus), Embeddings

Frameworks: React, Node.js, Next.js, Django, Flask

Developer Tools: Cursor, Jira, Confluence, MongoDB, Unix/Linux, Git